

Telecom Customer Churn Prediction Using Artificial Neural Network (ANN)

Group Members:

Member 1: Name: Rupali Patole Roll No.: MCS21005 Email Id: mcs21005@iiitl.ac.in

Member 2: Name: Aman Mehta Roll no: MCS21011 Email Id: mcs21011@iiitl.ac.in

Member 3: Name: Shreya Goswami Roll No.: MCS21023 Email Id: mcs21023@iiitl.ac.in

Introduction:

Customer churn prediction is to measure why customers are leaving a business. It can be applied in different business sectors like banking, retail, telecom, etc. For this project we will be implementing customer churn prediction for the telecom business. We will build a deep learning model to predict the churn and use precision, recall, f1-score to measure performance of our model.

Dataset: <https://www.kaggle.com/blastchar/telco-customer-churn>

Algorithm: Building ANN using tensorflow / keras

Steps:

1. Data Loading
 2. Data analysis
 3. Data cleaning
 4. Data splitting (Training set and Test set)
 5. ANN Model Building
 6. Evaluation Metrics (Accuracy, F1 score, Precision, Recall)
 7. Hyperparameter Tuning
 8. Performance Metrics (ROC Analysis)
 9. Conclusion
 10. Future Scope (Offers and Improving Churn Rate)
-

Code:

Imports

```
import numpy as np
import pandas as pd
import seaborn as sns
import plotly.express as px
import matplotlib.pyplot as plt
# %matplotlib inline
```

```

from sklearn import metrics

from sklearn.model_selection import train_test_split

import warnings

warnings.filterwarnings('ignore')

from keras.layers import Activation, Dense, Dropout

import tensorflow as tf

from keras.models import Sequential

from sklearn.ensemble import GradientBoostingClassifier

from sklearn.model_selection import GridSearchCV

from sklearn.metrics import plot_roc_curve, accuracy_score

```

Data Loading

```

df = pd.read_csv("customer_churn.csv")

df.sample(5)

```

Out[2]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection	TechS
6226	2097-YVPKN	Male	0	No	No	65	Yes	Yes	No	No internet service	...	No internet service	No i
4607	2853-CWQFQ	Male	0	No	Yes	1	Yes	No	DSL	No	...	No	
1010	0929-HYQEW	Male	0	No	No	3	Yes	No	DSL	Yes	...	No	
4932	5566-SOEZD	Male	0	Yes	Yes	27	Yes	No	Fiber optic	Yes	...	No	
6399	2101-RANCD	Female	0	No	No	55	Yes	Yes	Fiber optic	No	...	No	

5 rows × 21 columns

Data analysis and Data cleaning

Dropping customerID column as it is of no use

```

df.drop('customerID',axis='columns',inplace=True)

df.dtypes

```

```
Out[4]: gender          object
SeniorCitizen      int64
Partner            object
Dependents         object
tenure             int64
PhoneService       object
MultipleLines      object
InternetService    object
OnlineSecurity     object
OnlineBackup       object
DeviceProtection   object
TechSupport        object
StreamingTV        object
StreamingMovies    object
Contract           object
PaperlessBilling   object
PaymentMethod      object
MonthlyCharges     float64
TotalCharges       object
Churn              object
dtype: object
```

df.TotalCharges.values

```
Out[5]: array(['29.85', '1889.5', '108.15', ..., '346.45', '306.6', '6844.5'],
      dtype=object)
```

Finding out rows that are blank strings

pd.to_numeric(df.TotalCharges,errors='coerce').isnull()

```
Out[6]: 0      False
1      False
2      False
3      False
4      False
...
7038   False
7039   False
7040   False
7041   False
7042   False
Name: TotalCharges, Length: 7043, dtype: bool
```

df[pd.to_numeric(df.TotalCharges,errors='coerce').isnull()]

```
Out[7]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSu
488	Female	0	Yes	Yes	0	No	No phone service	DSL	Yes	No	Yes	
753	Male	0	No	Yes	0	Yes	No	No	No internet service	No internet service	No internet service	No in s
936	Female	0	Yes	Yes	0	Yes	No	DSL	Yes	Yes	Yes	
1082	Male	0	Yes	Yes	0	Yes	Yes	No	No internet service	No internet service	No internet service	No in s
1340	Female	0	Yes	Yes	0	No	No phone service	DSL	Yes	Yes	Yes	
3331	Male	0	Yes	Yes	0	Yes	No	No	No internet service	No internet service	No internet service	No in s
3826	Male	0	Yes	Yes	0	Yes	Yes	No	No internet service	No internet service	No internet service	No in s
4380	Female	0	Yes	Yes	0	Yes	No	No	No internet service	No internet service	No internet service	No in s
5218	Male	0	Yes	Yes	0	Yes	No	No	No internet service	No internet service	No internet service	No in s
6670	Female	0	Yes	Yes	0	Yes	Yes	DSL	No	Yes	Yes	
6754	Male	0	No	Yes	0	Yes	Yes	DSL	Yes	Yes	No	

df.shape

```
Out[8]: (7043, 20)
```

df.iloc[488].TotalCharges

```
Out[9]: ' '
```

df[df.TotalCharges!=' '].shape

```
Out[10]: (7032, 20)
```

Remove rows with space in TotalCharges

df1 = df[df.TotalCharges!=' ']

df1.shape

```
Out[11]: (7032, 20)
```

df1.dtypes

```
Out[12]: gender          object
SeniorCitizen         int64
Partner              object
Dependents            object
tenure                int64
PhoneService          object
MultipleLines         object
InternetService       object
OnlineSecurity        object
OnlineBackup          object
DeviceProtection      object
TechSupport           object
StreamingTV           object
StreamingMovies       object
Contract              object
PaperlessBilling      object
PaymentMethod         object
MonthlyCharges        float64
TotalCharges          object
Churn                 object
dtype: object
```

df1.TotalCharges = pd.to_numeric(df1.TotalCharges)

df1.TotalCharges.values

```
Out[14]: array([ 29.85, 1889.5 , 108.15, ..., 346.45, 306.6 , 6844.5 ])
```

df1[df1.Churn=='No']

Out[15]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSu
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	No	
1	Male	0	No	No	34	Yes	No	DSL	Yes	No	Yes	
3	Male	0	No	No	45	No	No phone service	DSL	Yes	No	Yes	
6	Male	0	No	Yes	22	Yes	Yes	Fiber optic	No	Yes	No	
7	Female	0	No	No	10	No	No phone service	DSL	Yes	No	No	
...
7037	Female	0	No	No	72	Yes	No	No	No internet service	No internet service	No internet service	No in s
7038	Male	0	Yes	Yes	24	Yes	Yes	DSL	Yes	No	Yes	
7039	Female	0	Yes	Yes	72	Yes	Yes	Fiber optic	No	Yes	Yes	
7040	Female	0	Yes	Yes	11	No	No phone service	DSL	Yes	No	No	
7042	Male	0	No	No	66	Yes	No	Fiber optic	Yes	No	Yes	

5163 rows x 20 columns

```
df1.Churn.value_counts()
```

```
Out[16]: No      5163
         Yes     1869
         Name: Churn, dtype: int64
```


```
plt.bar(x = df1['Churn'].unique(), height = df1.Churn.value_counts())
```

```
plt.legend()
```

```
plt.xlabel("Churn label")
```

```
plt.ylabel("Count")
```


```
plt.show()
```



```
plt.pie(df1.Churn.value_counts(), labels = df1['Churn'].unique(), explode = [0.1, 0.5])
```

```
plt.legend()
```

```
plt.show()
```



Data Visualization


```
tenure_churn_no = df1[df1.Churn=='No'].tenure  
tenure_churn_yes = df1[df1.Churn=='Yes'].tenure
```

```
plt.xlabel("tenure")  
plt.ylabel("Number Of Customers")  
plt.title("Customer Churn Prediction Visualiztion")
```

```
blood_sugar_men = [113, 85, 90, 150, 149, 88, 93, 115, 135, 80, 77, 82, 129]  
blood_sugar_women = [67, 98, 89, 120, 133, 150, 84, 69, 89, 79, 120, 112, 100]
```

```
plt.hist([tenure_churn_yes, tenure_churn_no], rwidth=0.95,  
color=['green','red'],label=['Churn=Yes','Churn=No'])  
plt.legend()
```

```
Out[19]: <matplotlib.legend.Legend at 0x212e377e5e0>
```




```
mc_churn_no = df1[df1.Churn=='No'].MonthlyCharges
mc_churn_yes = df1[df1.Churn=='Yes'].MonthlyCharges
```

```
plt.xlabel("Monthly Charges")
plt.ylabel("Number Of Customers")
plt.title("Customer Churn Prediction Visualiztion")
```

```
blood_sugar_men = [113, 85, 90, 150, 149, 88, 93, 115, 135, 80, 77, 82, 129]
blood_sugar_women = [67, 98, 89, 120, 133, 150, 84, 69, 89, 79, 120, 112, 100]
```

```
plt.hist([mc_churn_yes, mc_churn_no], rwidth=0.95, color=['green','red'],label=['Churn=Yes','Churn=No'])
plt.legend()
```

Out[20]: <matplotlib.legend.Legend at 0x2128a616fa0>



```
df1.groupby(['gender', 'Partner', 'Dependents', 'Churn']).size()
```

```
Out[21]: gender  Partner  Dependents  Churn
Female    No         No         No      1068
          No         No         Yes      587
          Yes        Yes        No      112
          Yes        Yes        Yes       33
          Yes        No         No      618
          Yes        Yes        Yes      187
          Yes        Yes        No      746
          Yes        Yes        Yes      132
Male      No         No         No     1089
          No         No         Yes      536
          Yes        No         No      170
          Yes        Yes        Yes       44
          Yes        No         No      615
          Yes        Yes        Yes      233
          Yes        Yes        No      745
          Yes        Yes        Yes      117
dtype: int64
```

```
df1.groupby(['gender', 'Partner', 'Dependents', 'Churn']).size().reset_index(name='Count')
```

Out[22]:


	gender	Partner	Dependents	Churn	Count
0	Female	No	No	No	1068
1	Female	No	No	Yes	587
2	Female	No	Yes	No	112
3	Female	No	Yes	Yes	33
4	Female	Yes	No	No	618
5	Female	Yes	No	Yes	187
6	Female	Yes	Yes	No	746
7	Female	Yes	Yes	Yes	132
8	Male	No	No	No	1089
9	Male	No	No	Yes	536
10	Male	No	Yes	No	170
11	Male	No	Yes	Yes	44
12	Male	Yes	No	No	615
13	Male	Yes	No	Yes	233
14	Male	Yes	Yes	No	745
15	Male	Yes	Yes	Yes	117

1. How gender, partner and dependents are related to churn?

```
fig = px.treemap(df1.groupby(['gender', 'Partner', 'Dependents', 'Churn']).size().reset_index(name='Count'),  
                 path=['gender', 'Partner', 'Dependents', 'Churn'], values='Count', color='Churn',  
                 title='1. How gender, partner and dependents are related to churn?')
```

fig.show()

1. How gender, partner and dependents are related to churn?



Observation: Whether male or female, if they do not have partner or dependents, they are more likely to churn! 2. Does tenure has any impact on churn?

```
df1.groupby(['tenure', 'Churn']).size().reset_index(name='count')
```

Out[24]:

	tenure	Churn	count
0	1	No	233
1	1	Yes	380
2	2	No	115
3	2	Yes	123
4	3	No	106
...
139	70	Yes	11
140	71	No	164
141	71	Yes	6
142	72	No	356
143	72	Yes	6


144 rows × 3 columns

2. Does tenure has any impact on churn?

```
fig = px.histogram(df1.groupby(['tenure', 'Churn']).size().reset_index(name='count'),
                   x="tenure", y="count", color="Churn", marginal="rug", color_discrete_map={"Yes": "#E45756",
                                                                                          "No": "#1CBE4F"},
                   title="2. Does tenure has any impact on churn?")
```

```
fig.show()
```

2. Does tenure has any impact on churn?



Observation: During 0-10 years of tenure, we can see maximum churning. As the customer turns old, they might get habituated using same telecom service

3. As the dataset is about telecom industry, we need some insights on phone and internet services!

```
df1.groupby(['Churn', 'PhoneService', 'InternetService']).size()
```

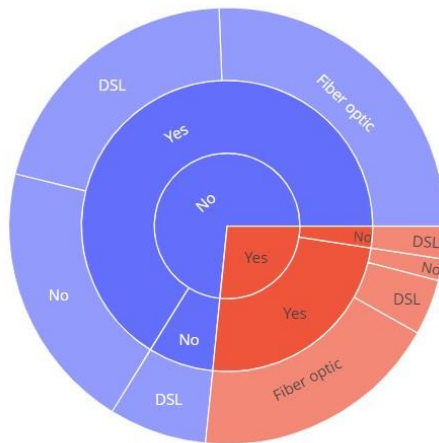
```
Out[26]: Churn  PhoneService  InternetService
No      No      DSL             510
        Yes     DSL             1447
        Yes     Fiber optic     1799
Yes     No      DSL             170
        Yes     DSL             289
        Yes     Fiber optic     1297
        Yes     No             113
dtype: int64
```

3. As the dataset is about telecom industry, we need some insights on phone and internet services!

```
fig = px.sunburst(df1.groupby(['Churn', 'PhoneService', 'InternetService']).size().reset_index(name='count'),
                  path = ['Churn', 'PhoneService', 'InternetService'], values = 'count',
                  title='3. As the dataset is about telecom industry, we need some insights on phone and internet
services!')
```

```
fig.show()
```

3. As the dataset is about telecom industry, we need some insights on phone and internet services!



```
np.unique(df1.TechSupport)
```

```
Out[28]: array(['No', 'No internet service', 'Yes'], dtype=object)
```

```
data_techSupport_yes = df1[df1['TechSupport'] == 'Yes']
```

```
data_techSupport_no = df1[df1['TechSupport'] == 'No']
```

Customers who took tech support

```
data_techSupport_yes.groupby(['tenure', 'Churn']).size()
```

```
Out[30]:
```

tenure	Churn	
1	No	19
	Yes	14
2	No	8
	Yes	11
3	No	20
	...	
70	Yes	8
71	No	89
	Yes	2
72	No	221
	Yes	3

Length: 140, dtype: int64

```
fig = px.histogram(data_techSupport_yes.groupby(['tenure', 'Churn']).size().reset_index(name='count'),  
                   x="tenure", y="count", marginal="rug", color="Churn", color_discrete_map={"Yes": "#E45756",  
                   "No": "#1CBE4F"},  
                   title="Statistics of customers opted for tech support with churning")
```

```
fig.show()
```

Statistics of customers opted for tech support with churning



Customers who didn't took tech support

```
data_techSupport_no.groupby(['tenure', 'Churn']).size()
```

```
Out[32]:
```

tenure	Churn	count
1	No	106
	Yes	308
2	No	66
	Yes	106
3	No	53
	...	
70	Yes	3
71	No	34
	Yes	4
72	No	69
	Yes	3

Length: 144, dtype: int64

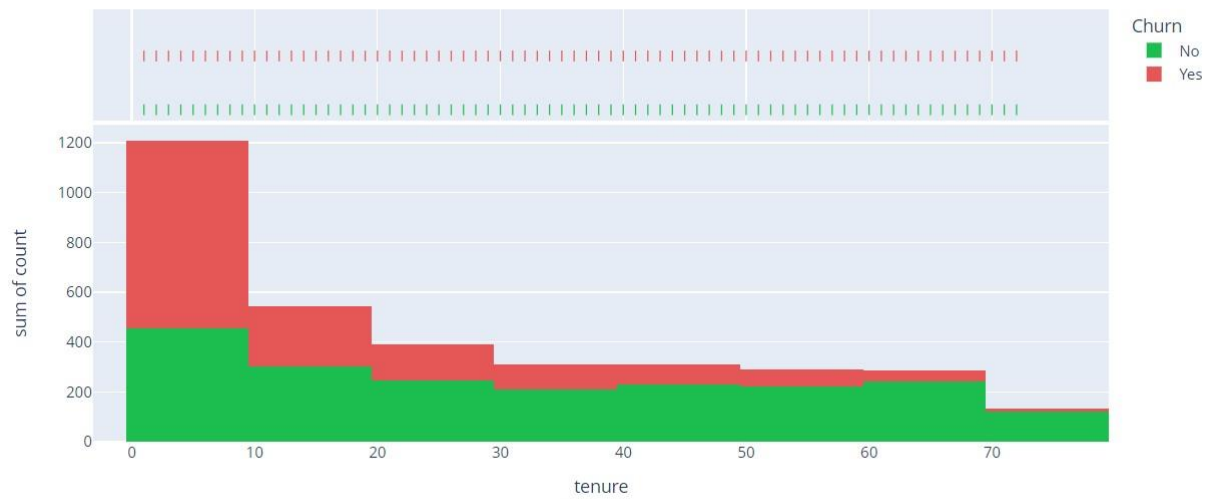
```
fig = px.histogram(data_techSupport_no.groupby(['tenure', 'Churn']).size().reset_index(name='count'),
```

```
                    x='tenure', y='count', color='Churn', marginal='rug', color_discrete_map={"Yes": "#E45756",  
"No": "#1CBE4F"},
```

```
                    title="Statistics of customers opted for tech support with churning")
```

```
fig.show()
```

Statistics of customers opted for tech support with churning



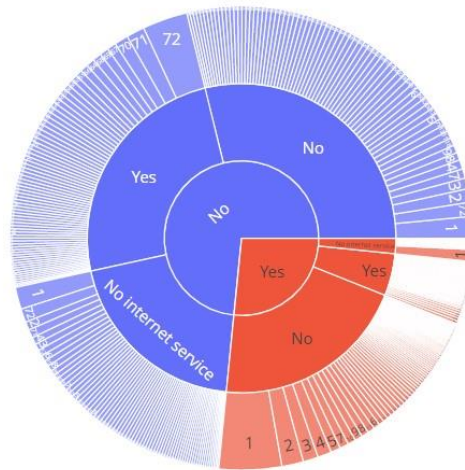
Observations: People with Phone services (yes) and 'Fiber optic' Internet Service are churning more

4. Does customers opted for tech support stayed for longer tenure with less churn?

```
fig = px.sunburst(df1.groupby(['Churn', 'TechSupport', 'tenure']).size().reset_index(name='count'),  
                  path=['Churn', 'TechSupport', 'tenure'], values='count',  
                  title='4. Does customers opted for tech support stayed for longer tenure with less churn?')
```

```
fig.show()
```

4. Does customers opted for tech support stayed for longer tenure with less churn?



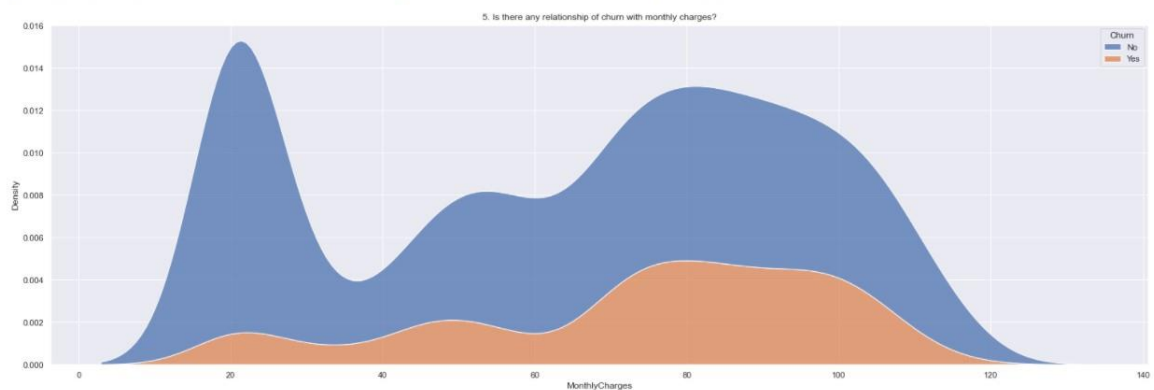
Observations: more churning takes place in first 10 yrs (max in first year itself), for customers with or without tech support. But Churning is more in case of "without tech support" customers

5. Is there any relationship of churn with monthly charges or total charges?

```
sns.set(rc={'figure.figsize':(26, 8.27)}) # rc - seems row, column
```

```
sns.kdeplot(data = df1, x="MonthlyCharges", hue="Churn", multiple="stack").set(title= "5. Is there any relationship of churn with monthly charges?")
```

```
Out[35]: [Text(0.5, 1.0, '5. Is there any relationship of churn with monthly charges?')]
```

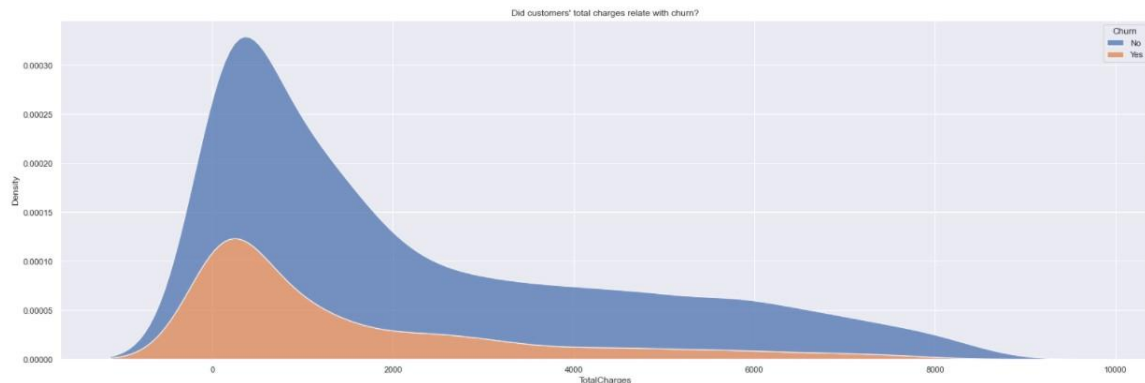


Observations: As the monthly charges are increasing, we can see the density increasing too (60-120), which means more churning with increasing monthly charges

```
sns.set(rc={'figure.figsize':(26,8.27)})
```

```
sns.kdeplot(data=df1, x="TotalCharges", hue="Churn", multiple="stack").set(title="Did customers' total charges relate with churn?")
```

```
Out[36]: [Text(0.5, 1.0, "Did customers' total charges relate with churn?")]
```



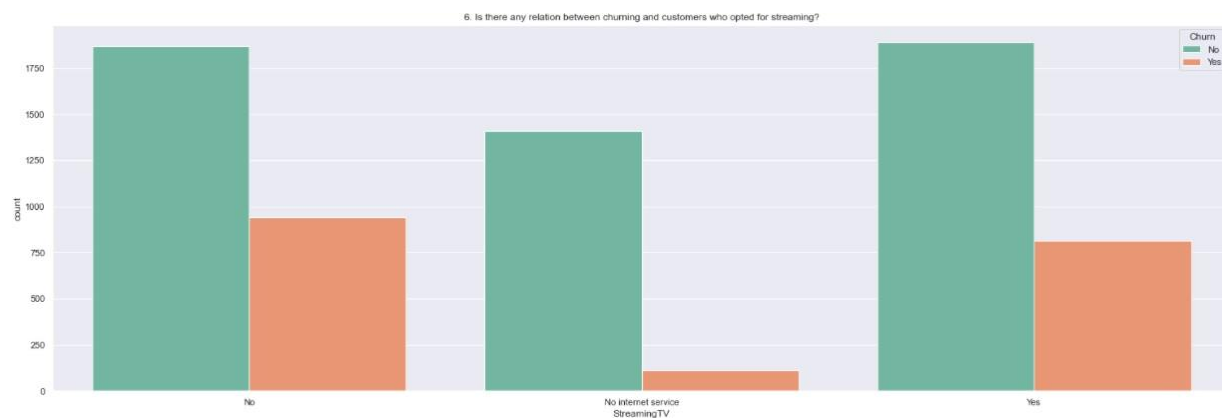
Observation: It is quite opposite of what has been seen for monthly charges. Here high churning occurs when total charges is less, 0-2000 total charges have maximum churning

```
df1.groupby(['Churn', 'StreamingTV']).size()
```

```
Out[37]: Churn  StreamingTV
No         No           1867
          No internet service 1407
          Yes           1889
Yes        No           942
          No internet service 113
          Yes            814
dtype: int64
```

```
ax = sns.barplot(x="StreamingTV", y="count", hue="Churn", data = df1.groupby(['Churn', 'StreamingTV']).size().reset_index(name='count'),
```

```
palette="Set2").set(title = "6. Is there any relation between churning and customers who opted for streaming?")
```



6. Is there any relation between churning and customers who opted for streaming?

```
fig = px.sunburst(df1.groupby(['Churn', 'InternetService', 'StreamingTV']).size().reset_index(name='count'),
                  path=['Churn', 'InternetService', 'StreamingTV'], values='count',
                  title='6. Is there any relation between churning and customers who opted for streaming?')
```

```
fig.show()
```


6. Is there any relation between churning and customers who opted for streaming?



```
fig = px.sunburst(df1.groupby(['Churn', 'StreamingTV']).size().reset_index(name='count'),
                  path=['Churn', 'StreamingTV'], values='count',
                  title='Do customers opted for streaming, faced issue with the service?')
```

```
fig.show()
```


Do customers opted for streaming, faced issue with the service?




Observation: Churning is being observed equally for the 'Yes', 'No' group of whether connected StreamingTv or not!

```
df1.groupby(['Churn', 'StreamingMovies']).size()
```

```
Out[41]: Churn    StreamingMovies
No         No          1843
         No internet service 1407
         Yes           1913
Yes        No           938
         No internet service 113
         Yes            818
dtype: int64
```

```
ax = sns.barplot(x="Churn", y="count", hue="StreamingMovies",
                 data = df1.groupby(['Churn', 'StreamingMovies']).size().reset_index(name="count"),
                 palette="Set2").set(title="")
```



Observation: Churning is being observed equally for both the 'Yes', 'No' group of StreamingMovies

```
df1.groupby(['Churn', 'Contract']).size()
```

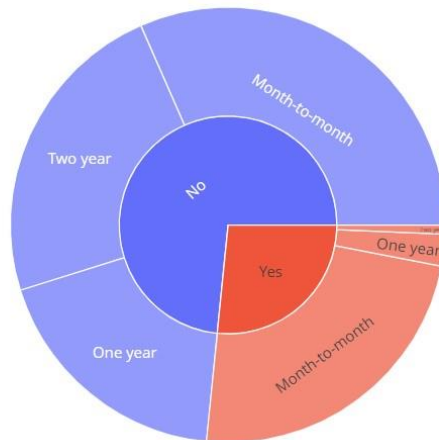
```
Out[43]: Churn  Contract      count
        No    Month-to-month    2220
         One year    1306
         Two year    1637
        Yes    Month-to-month    1655
         One year    166
         Two year     48
dtype: int64
```

7. How contract is impacting business?

```
fig = px.sunburst(df1.groupby(['Churn', 'Contract']).size().reset_index(name='count'),
                  path=['Churn', 'Contract'], values='count',
                  title='7. How contract is impacting business?')
```

```
fig.show()
```

7. How contract is impacting business?




Observations: clearly visible that customers with month-to-month contract are the highest churners

Senior Citizen vs Churning

```
fig = px.sunburst(df1.groupby(['Churn', 'SeniorCitizen']).size().reset_index(name='count'),
                  path=['Churn', 'SeniorCitizen'], values='count',
                  title='How being or non being SeniorCitizen is impacting Churning?')
```

fig.show()

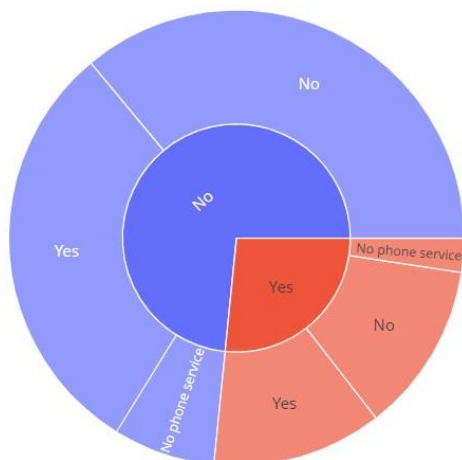
How being or non being SeniorCitizen is impacting Churning?



```
fig = px.sunburst(df1.groupby(['Churn', 'MultipleLines']).size().reset_index(name='count'),  
                 path=['Churn', 'MultipleLines'], values='count',  
                 title='How having MultipleLines is impacting Churning?')
```

fig.show()

How having MultipleLines is impacting Churning?



Observation: Having (yes) multiple lines have almost equal impact as not having (No) multiple lines

Many of the columns are yes, no etc. Let's print unique values in object columns to see data values

```
def print_unique_col_values(df):  
    for column in df:  
        if df[column].dtypes=='object':  
            print(f'{column}: {df[column].unique()}')  
  
print_unique_col_values(df1)  
  
gender: ['Female' 'Male']  
Partner: ['Yes' 'No']  
Dependents: ['No' 'Yes']  
PhoneService: ['No' 'Yes']  
MultipleLines: ['No phone service' 'No' 'Yes']  
InternetService: ['DSL' 'Fiber optic' 'No']  
OnlineSecurity: ['No' 'Yes' 'No internet service']  
OnlineBackup: ['Yes' 'No' 'No internet service']  
DeviceProtection: ['No' 'Yes' 'No internet service']  
TechSupport: ['No' 'Yes' 'No internet service']  
StreamingTV: ['No' 'Yes' 'No internet service']  
StreamingMovies: ['No' 'Yes' 'No internet service']  
Contract: ['Month-to-month' 'One year' 'Two year']  
PaperlessBilling: ['Yes' 'No']  
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'  
                'Credit card (automatic)']  
Churn: ['No' 'Yes']
```

Some of the columns have no internet service or no phone service, that can be replaced with a simple No

```
df1.replace('No internet service','No',inplace=True)
```

```
df1.replace('No phone service','No',inplace=True)
```

```
print_unique_col_values(df1)  
  
gender: ['Female' 'Male']  
Partner: ['Yes' 'No']  
Dependents: ['No' 'Yes']  
PhoneService: ['No' 'Yes']  
MultipleLines: ['No' 'Yes']  
InternetService: ['DSL' 'Fiber optic' 'No']  
OnlineSecurity: ['No' 'Yes']  
OnlineBackup: ['Yes' 'No']  
DeviceProtection: ['No' 'Yes']  
TechSupport: ['No' 'Yes']  
StreamingTV: ['No' 'Yes']  
StreamingMovies: ['No' 'Yes']  
Contract: ['Month-to-month' 'One year' 'Two year']  
PaperlessBilling: ['Yes' 'No']  
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'  
                'Credit card (automatic)']  
Churn: ['No' 'Yes']
```

Convert Yes and No to 1 or 0

```
yes_no_columns = ['Partner','Dependents','PhoneService','MultipleLines','OnlineSecurity','OnlineBackup',
                  'DeviceProtection','TechSupport','StreamingTV','StreamingMovies','PaperlessBilling','Churn']
```

for col in yes_no_columns:

```
    df1[col].replace({'Yes': 1, 'No': 0}, inplace=True)
```

for col in df1:

```
    print(f'{col}: {df1[col].unique()}')
```

```
gender: ['Female' 'Male']
SeniorCitizen: [0 1]
Partner: [1 0]
Dependents: [0 1]
tenure: [ 1 34  2 45  8 22 10 28 62 13 16 58 49 25 69 52 71 21 12 30 47 72 17 27
         5 46 11 70 63 43 15 60 18 66  9  3 31 50 64 56  7 42 35 48 29 65 38 68
        32 55 37 36 41  6  4 33 67 23 57 61 14 20 53 40 59 24 44 19 54 51 26 39]
PhoneService: [0 1]
MultipleLines: [0 1]
InternetService: ['DSL' 'Fiber optic' 'No']
OnlineSecurity: [0 1]
OnlineBackup: [1 0]
DeviceProtection: [0 1]
TechSupport: [0 1]
StreamingTV: [0 1]
StreamingMovies: [0 1]
Contract: ['Month-to-month' 'One year' 'Two year']
PaperlessBilling: [1 0]
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
                'Credit card (automatic)']
MonthlyCharges: [29.85 56.95 53.85 ... 63.1 44.2 78.7 ]
TotalCharges: [ 29.85 1889.5 108.15 ... 346.45 306.6 6844.5 ]
Churn: [0 1]
```

```
df1['gender'].replace({'Female':1,'Male':0},inplace=True)
```

```
df1.gender.unique()
```

```
Out[54]: array([1, 0], dtype=int64)
```

One hot encoding for categorical columns

```
df2 = pd.get_dummies(data=df1, columns=['InternetService','Contract','PaymentMethod'])
```

```
df2.columns
```

```
Out[55]: Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
               'PhoneService', 'MultipleLines', 'OnlineSecurity', 'OnlineBackup',
               'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
               'PaperlessBilling', 'MonthlyCharges', 'TotalCharges', 'Churn',
               'InternetService_DSL', 'InternetService_Fiber optic',
               'InternetService_No', 'Contract_Month-to-month', 'Contract_One year',
               'Contract_Two year', 'PaymentMethod_Bank transfer (automatic)',
               'PaymentMethod_Credit card (automatic)',
               'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check'],
              dtype='object')
```

df2.sample(5)

```
Out[56]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	OnlineSecurity	OnlineBackup	DeviceProtection	...	InternetService_DSL
2814	1	0	1	1	62	0	0	1	1	0	...	1
4933	1	0	0	1	4	0	0	1	1	1	...	1
2040	0	0	1	0	71	1	1	1	1	1	...	1
4772	1	0	1	1	69	1	0	0	1	1	...	1
1674	1	0	1	1	23	1	1	0	0	0	...	0

5 rows × 27 columns

df2.dtypes

```
Out[57]: gender                int64
SeniorCitizen                int64
Partner                      int64
Dependents                   int64
tenure                       int64
PhoneService                 int64
MultipleLines                int64
OnlineSecurity               int64
OnlineBackup                 int64
DeviceProtection             int64
TechSupport                  int64
StreamingTV                  int64
StreamingMovies              int64
PaperlessBilling              int64
MonthlyCharges                float64
TotalCharges                  float64
Churn                        int64
InternetService_DSL           uint8
InternetService_Fiber optic   uint8
InternetService_No            uint8
Contract_Month-to-month       uint8
Contract_One year             uint8
Contract_Two year             uint8
PaymentMethod_Bank transfer (automatic) uint8
PaymentMethod_Credit card (automatic)  uint8
PaymentMethod_Electronic check         uint8
PaymentMethod_Mailed check             uint8
dtype: object
```

```
cols_to_scale = ['tenure','MonthlyCharges','TotalCharges']
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
```

```
df2[cols_to_scale] = scaler.fit_transform(df2[cols_to_scale])
```

```
for col in df2:
```

```

print(f'{col}: {df2[col].unique()}')

gender: [1 0]
SeniorCitizen: [0 1]
Partner: [1 0]
Dependents: [0 1]
tenure: [0. 0.46478873 0.01408451 0.61971831 0.09859155 0.29577465
0.12676056 0.38028169 0.85915493 0.16901408 0.21126761 0.8028169
0.67605634 0.33802817 0.95774648 0.71830986 0.98591549 0.28169014
0.15492958 0.4084507 0.64788732 1. 0.22535211 0.36619718
0.05633803 0.63380282 0.14084507 0.97183099 0.87323944 0.5915493
0.1971831 0.83098592 0.23943662 0.91549296 0.11267606 0.02816901
0.42253521 0.69014085 0.88732394 0.77464789 0.08450704 0.57746479
0.47887324 0.66197183 0.3943662 0.90140845 0.52112676 0.94366197
0.43661972 0.76056338 0.50704225 0.49295775 0.56338028 0.07042254
0.04225352 0.45070423 0.92957746 0.30985915 0.78873239 0.84507042
0.18309859 0.26760563 0.73239437 0.54929577 0.81690141 0.32394366
0.6056338 0.25352113 0.74647887 0.70422535 0.35211268 0.53521127]
PhoneService: [0 1]
MultipleLines: [0 1]
OnlineSecurity: [0 1]
OnlineBackup: [1 0]
DeviceProtection: [0 1]
TechSupport: [0 1]
StreamingTV: [0 1]
StreamingMovies: [0 1]
PaperlessBilling: [1 0]
MonthlyCharges: [0.11542289 0.38507463 0.35422886 ... 0.44626866 0.25820896 0.60149254]
TotalCharges: [0.0012751 0.21586661 0.01031041 ... 0.03780868 0.03321025 0.78764136]
Churn: [0 1]
InternetService_DSL: [1 0]
InternetService_Fiber optic: [0 1]
InternetService_No: [0 1]
Contract Month-to-month: [1 0]

```

Data Splitting

Train Test Split

```

X = df2.drop('Churn',axis='columns')
y = df2['Churn']

X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=5)

```

X_train.shape

```
Out[61]: (5625, 26)
```

X_test.shape

```
Out[62]: (1407, 26)
```

X_train[:10]

Out[63]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	OnlineSecurity	OnlineBackup	DeviceProtection	...	InternetService_D!
5664	1	1	0	0	0.126761	1	0	0	0	1	...	
101	1	0	1	1	0.000000	1	0	0	0	0	...	
2621	0	0	1	0	0.985915	1	0	0	1	1	...	
392	1	1	0	0	0.014085	1	0	0	0	0	...	
1327	0	0	1	0	0.816901	1	1	0	0	1	...	
3607	1	0	0	0	0.169014	1	0	1	0	0	...	
2773	0	0	1	0	0.323944	0	0	0	0	1	...	
1936	1	0	1	0	0.704225	1	0	1	1	0	...	
5387	0	0	0	0	0.042254	0	0	0	0	0	...	
4331	0	0	0	0	0.985915	1	1	0	0	0	...	

10 rows × 26 columns

len(X_train.columns)

Out[64]: 26

ANN Model Building

```
nn_model = Sequential()
```

```
nn_model.add(Dense(64, kernel_regularizer=tf.keras.regularizers.l2(0.001), input_dim=26, activation='relu'))
```

```
nn_model.add(Dropout(rate=0.2))
```

```
nn_model.add(Dense(8, kernel_regularizer=tf.keras.regularizers.l2(0.001), activation='relu'))
```

```
nn_model.add(Dropout(rate=0.1))
```

```
nn_model.add(Dense(1, activation='sigmoid'))
```

```
lr_schedule = tf.keras.optimizers.schedules.InverseTimeDecay( 0.001,
```

```
    decay_steps=(X_train.shape[0]/32)*50,
```

```
    decay_rate=1,
```

```
    staircase=False)
```

This time decay means for every 50 epochs the learning rate will be half of 0.001 value

```
def get_optimizer():
```

```
    return tf.keras.optimizers.Adam(lr_schedule)
```

```
def get_callbacks():
```


```
    return [tf.keras.callbacks.EarlyStopping(monitor='val_accuracy',patience=70,restore_best_weights=True)]
```



```
nn_model.compile(loss = "binary_crossentropy",
                 optimizer = get_optimizer(),
                 metrics=['accuracy'])
```

```
history = nn_model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=150, batch_size=32,
                      callbacks=get_callbacks(), verbose=0)
```

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



Evaluation Metrics

```
yprednn=nn_model.predict(X_test)
yprednn=yprednn.round()
print('Neural Network:\n {} \n'.format(
    metrics.classification_report(yprednn, y_test)))
nn_conf_matrix=metrics.confusion_matrix(yprednn,y_test)
conf_mat_nn = pd.DataFrame(nn_conf_matrix,
```

```

columns=["Predicted NO", "Predicted YES"],
index=["Actual NO", "Actual YES"])
print(conf_mat_nn)

```

Accuracy with ANN Model: 80.0%

```

Neural Network:
      precision    recall  f1-score   support

    0.0         0.89     0.83     0.86     1067
    1.0         0.57     0.68     0.62      340

 accuracy
macro avg         0.73     0.76     0.74     1407
weighted avg         0.81     0.80     0.80     1407

```

```

      Predicted NO Predicted YES
Actual NO         890         177
Actual YES         109         231

```

Accuracy with ANN Model: 80.0%

Hyperparameter Tuning

```
gb = GradientBoostingClassifier()
```

```

parameters = {
    'loss': ['deviance', 'exponential'],
    'learning_rate': [0.001, 0.1, 1, 10],
    'n_estimators': [100, 150, 180, 200]
}

```

```
grid_search = GridSearchCV(gb, parameters, cv = 5, n_jobs = -1, verbose = 1)
```

```
grid_search.fit(X_train, y_train)
```

```

Fitting 5 folds for each of 32 candidates, totalling 160 fits
Out[76]: GridSearchCV(cv=5, estimator=GradientBoostingClassifier(), n_jobs=-1,
      param_grid={'learning_rate': [0.001, 0.1, 1, 10],
      'loss': ['deviance', 'exponential'],
      'n_estimators': [100, 150, 180, 200]},
      verbose=1)

```

```
print(grid_search.best_params_)
```

```
print(grid_search.best_score_)
```

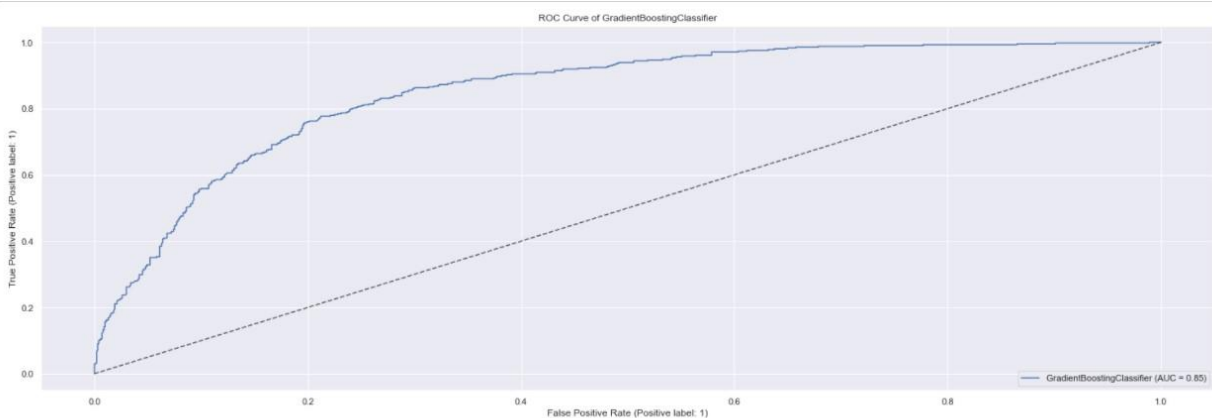
```
{'learning_rate': 0.1, 'loss': 'exponential', 'n_estimators': 100}  
0.8071111111111111
```

```
gb = GradientBoostingClassifier(learning_rate = 0.1, loss = 'exponential', n_estimators = 80)  
gb.fit(X_train,y_train)  
y_pred = gb.predict(X_test)  
print("Accuracy with Gradient Booting Classifier: {0:.5}%".format(accuracy_score(y_test, y_pred)*100))
```

```
Accuracy with Gradient Booting Classifier: 80.028%
```

Performance Metrics

```
gbc_disp = plot_roc_curve(gb, X_test, y_test)  
plt.title("ROC Curve of {}".format(type(gb).__name__))  
plt.plot([0,1],[0,1],"--",color="k",alpha=0.7)  
plt.show()
```



Conclusion:

1. When we built a Model with Artificial Neural Network using Tensorflow/Keras, we found that the accuracy was 80%.

2. Hyperparameter Tuning using Gradient Boosting Classifier yielded an accuracy score of 80.028% and an AUC value of 0.85.

Thus, we can conclude that our ANN Model has performed appropriately.

Future Scope:

Offers and Improving Churn Rate:

1. Discounts: As the most important feature is total charges, followed by monthly charges, potential churners identified through the modelling should be offered huge discount on next month or months contract. This covers 80 % of the reasons identified for churning. For this modelling, the False Negative Rate should be minimized or Recall should be maximized so that the discounts are sent to maximum of the potential churners.

2. New contract: A six month or four month contract should be implemented. This will encourage the reserved customers who want shorter contracts and will increase their tenure on the service thus making them less likely to churn.

3. Online Security: This service should be promoted more and offered complimentary/free for trial periods depending on cost to company. The customers who do not have online security are more likely to churn and thus this offer could be combined with the first one mentioned and discount could only be offered on this.

4. Fiber optic: The fiber optic internet is costly and thus should either be promoted to appropriate target audience or better technology can be implemented to cut cost on this service. Ultimately the market research team has to decide the break even point for this service, whether it is profiting as much as the loss of customers it is causing.

5. Another method to quantify the offers to be made is using manually generated features and their effect on the model.
